

Lecture 1:

Math (P)Review Part I: Linear Algebra

Computer Graphics

CMU 15-462/15-662, Fall 2016

Homework 0.5 (Out later today!)

- Exercises will generally be a bit harder / more rigorous than what you will do for the rest of the class.
- Goal is to help you build strength for the upcoming journey.
- We are here to help!

1 Linear Algebra

1.1 Basic Vector Operations

Exercise 1. Letting $\mathbf{u} := (4, 3)$, $\mathbf{v} := (4, 3)$, $a := 7$ and $b := 7$, calculate the following quantities:

- (a) $\mathbf{u} + \mathbf{v}$
- (b) $b\mathbf{u}$
- (c) $a\mathbf{u} - b\mathbf{v}$

Exercise 2. Letting $\mathbf{u} := (8, 2, 7)$ and $\mathbf{v} := (8, 7, 3)$, calculate the following quantities:


- 1. $\mathbf{u} - \mathbf{v}$
- 2. $\mathbf{u} + 6\mathbf{v}$

Exercise 3. So far we have been working with vectors in \mathbb{R}^2 and \mathbb{R}^3 , but it is important to remember that other objects, like functions, also behave like vectors in the sense that we can add them, subtract them, multiply them by scalars, etc. Calculate the following quantities for the two polynomials $p(x) := 8x^2 + 2x + 7$ and $q(x) := 8x^2 + 7x + 3$, and evaluate the result at the point $x = 7$:

- 1. $p(x) - q(x)$
- 2. $p(x) + 6q(x)$


Linear Algebra in Computer Graphics

- Today's topic: **linear algebra**.
- Why is linear algebra important for computer graphics?
 - Effective bridge between geometry, physics, etc., and *computation*.
 - In many areas of graphics, once you can express the solution to a problem in terms of linear algebra, you're essentially done: now ask the computer to solve $Ax=b$.
 - Development of fast numerical linear algebra has made modern computer graphics possible (image processing, physically-based animation, geometry processing...)



Vectors - Intuition

- Linear algebra is the study of **vector spaces** and **linear maps** between them.
- First things first: what is a **vector**?
- Intuitively, a vector is a little arrow:




A vector.

- In computer graphics, we work with many types of data that may not look like little arrows (polynomials, images, radiance...). But they still *behave* like vectors. So, this little arrow is still often a useful mental model.

Vectors - What Can We Measure?

- What information does a vector encode?
- Fundamentally, just **direction** and **magnitude***:



- For instance, a vector in 2D can be encoded by a length and an angle relative to some fixed direction (“polar coordinates”).
- How else might we encode a vector?


Traditionally, a vector does not include a “basepoint”; a vector with a basepoint is sometimes called a **tangent vector.*

Vector in Cartesian Coordinates

- Can also measure components of a vector with respect to some chosen *coordinate system*:




René Descartes, Est. 1596



- **WARNING:** Can't directly compare coordinates in different systems! (Also shouldn't compare (r, θ) to (x, y) .)

What Can We Do with a Vector?

- Two basic operations. First, we can add them “end to end”:




- What if we walk along v first, then u ?
- Actually, it doesn't seem to matter: $u + v = v + u$
- Language: vector addition is “commutative” or “abelian”



Niels Henrik Abel

What Else Can We Do with a Vector?

- Other basic operation? Scaling:




- In general, can multiply any vector u by a number or “scalar” a to get a new vector au .
- Multiplication behaves the way we would expect, based on the geometric behavior of scaling “little arrows.” E.g.,

$$a(bu) = (ab)u$$

Interaction of Addition & Scaling

- What if we try to add two scaled vectors? Or scale two vectors that have been added together?



- Interesting—seems we get the same result either way:

$$a(\mathbf{u} + \mathbf{v}) = a\mathbf{u} + a\mathbf{v}$$

Vector Space—Formal Definition

- If we keep playing around vectors, eventually we come up with a complete set of “rules” that vectors seem to obey*:

For all vectors $\mathbf{u}, \mathbf{v}, \mathbf{w}$ and scalars a, b :

- $\mathbf{u} + \mathbf{v} = \mathbf{v} + \mathbf{u}$
- $\mathbf{u} + (\mathbf{v} + \mathbf{w}) = (\mathbf{u} + \mathbf{v}) + \mathbf{w}$
- There exists a *zero vector* “ $\mathbf{0}$ ” such that $\mathbf{v} + \mathbf{0} = \mathbf{0} + \mathbf{v} = \mathbf{v}$
- For every \mathbf{v} there is a vector “ $-\mathbf{v}$ ” such that $\mathbf{v} + (-\mathbf{v}) = \mathbf{0}$
- $1\mathbf{v} = \mathbf{v}$
- $a(b\mathbf{v}) = (ab)\mathbf{v}$
- $a(\mathbf{u} + \mathbf{v}) = a\mathbf{u} + a\mathbf{v}$
- $(a + b)\mathbf{v} = a\mathbf{v} + b\mathbf{v}$

- ***These rules did not “fall out of the sky!”*** Each one comes from the geometric behavior of “little arrows.” (Can you draw a picture for each one?)
- ***Any collection of objects satisfying all of this properties is a **vector space** (even if they don’t look like little arrows!)***

***this will NOT be on the test!**

Euclidean Vector Space

- **Most common example: Euclidean n-dimensional space**
- **Typically denoted by \mathbb{R}^n , meaning “n real numbers”**
- **E.g., $(1.23, 4.56, \pi/2)$ is a point in \mathbb{R}^3**
- **Why such a common example?**
 - **Looks a lot like the space we live in!**
 - **That’s what we can easily encode on a computer (a list of floating-point numbers).**

\mathbb{R}



\mathbb{R}^2











\mathbb{R}^3




Functions as Vectors

- Another very important example of vector spaces in computer graphics are spaces of *functions*.
- Why? Because many of the objects we want to work with in graphics are functions! (Images, radiance from a light source, surfaces, modal vibrations, ...)




Vector Operations on Functions

- Do functions exhibit the same behavior as “little arrows?”
- Well, we can certainly add two functions:



- We can also scale a function:



Vector Operations on Functions

■ What about the rest of these properties?

For all vectors $\mathbf{u}, \mathbf{v}, \mathbf{w}$ and scalars a, b :

- $\mathbf{u} + \mathbf{v} = \mathbf{v} + \mathbf{u}$
- $\mathbf{u} + (\mathbf{v} + \mathbf{w}) = (\mathbf{u} + \mathbf{v}) + \mathbf{w}$
- There exists a *zero vector* “ $\mathbf{0}$ ” such that $\mathbf{v} + \mathbf{0} = \mathbf{0} + \mathbf{v} = \mathbf{v}$
- For every \mathbf{v} there is a vector “ $-\mathbf{v}$ ” such that $\mathbf{v} + (-\mathbf{v}) = \mathbf{0}$
- $1\mathbf{v} = \mathbf{v}$
- $a(b\mathbf{v}) = (ab)\mathbf{v}$
- $a(\mathbf{u} + \mathbf{v}) = a\mathbf{u} + b\mathbf{v}$
- $(a + b)\mathbf{v} = a\mathbf{v} + b\mathbf{v}$


■ Try it out at home!

■ E.g., the “zero vector” is the function equal to zero for all x .

■ Short answer: yes, functions are vectors! (Even if they don’t look like “little arrows”.)

Manipulating Vectors in Coordinates

- So far, we've only drawn our vector operations via pictures.
- How do we actually compute with vectors?
- Return to our coordinate representation:



$$\begin{aligned}\mathbf{u} + \mathbf{v} &= (1, 3) + (4, 1) \\ &= (1 + 4, 3 + 1) \\ &= (5, 4)\end{aligned}$$

**Ok, so we came up with some
rule for adding pairs of numbers.**

**How can we check that it faithfully encodes
geometric behavior of “little arrows?”**

From Geometry to Algebra

- Just check that it agrees with our list of rules that we know (from reasoning *geometrically*) “little arrows” must obey:

For all vectors $\mathbf{u}, \mathbf{v}, \mathbf{w}$ and scalars a, b :

- $\mathbf{u} + \mathbf{v} = \mathbf{v} + \mathbf{u}$
- $\mathbf{u} + (\mathbf{v} + \mathbf{w}) = (\mathbf{u} + \mathbf{v}) + \mathbf{w}$
- There exists a *zero vector* “ $\mathbf{0}$ ” such that $\mathbf{v} + \mathbf{0} = \mathbf{0} + \mathbf{v} = \mathbf{v}$
- For every \mathbf{v} there is a vector “ $-\mathbf{v}$ ” such that $\mathbf{v} + (-\mathbf{v}) = \mathbf{0}$
- $1\mathbf{v} = \mathbf{v}$
- $a(b\mathbf{v}) = (ab)\mathbf{v}$
- $a(\mathbf{u} + \mathbf{v}) = a\mathbf{u} + a\mathbf{v}$
- $(a + b)\mathbf{v} = a\mathbf{v} + b\mathbf{v}$

- For instance, for any two vectors $\mathbf{u} := (u_1, u_2)$ and $\mathbf{v} := (v_1, v_2)$ we have

$$\begin{aligned}\mathbf{u} + \mathbf{v} &= (u_1, u_2) + (v_1, v_2) = (u_1 + v_1, u_2 + v_2) = \\ &= (v_1 + u_1, v_2 + u_2) = (v_1, v_2) + (u_1, u_2) = \mathbf{v} + \mathbf{u}.\end{aligned}$$


Turning geometric observations into algebraic rules is convenient for symbolic manipulation & numerical computation.

But you should *never* blindly accept a rule given purely by authority.

**Always ask: where does this rule come from?
What does it mean geometrically?
(Can you draw a picture?)**

Scaling Vectors in Coordinates

- We'd also like to be able to scale vectors using coordinates.
- Any ideas?




$$\begin{aligned}\frac{3}{2}\mathbf{u} \\ &= \frac{3}{2}(4, 2) \\ &= (4 \cdot 3/2, 2 \cdot 3/2) \\ &= (12/2, 6/2) \\ &= (6, 3)\end{aligned}$$

(From here, check the rest of the properties...)

Computing the Midpoint


- As we start to combine vector operations, we build up operations needed for computer graphics.
- E.g., how would I compute the midpoint m of $a = (3,4)$ and $b = (7,2)$?



$$\begin{aligned} \mathbf{m} &= \frac{1}{2}(\mathbf{a} + \mathbf{b}) \\ &= \frac{1}{2}((3,4) + (7,2)) \\ &= \frac{1}{2}(10,6) \\ &= (5,3) \end{aligned}$$


Measuring Vectors

- Earlier we asked, “what information does a vector encode?”
- (A: Orientation and magnitude.)
- How do we actually *measure* these quantities?



Norm of a Vector

- Let's start with magnitude—for a given vector v , we want to assign it a number $|v|$ called its **length** or **magnitude** or **norm**.
- Intuitively, the norm should capture how “big” the vector is.



small norm



large norm

Natural Properties of Length—Positivity

- What properties might you expect the norm (or length) of a vector to satisfy?
- For one thing, it probably shouldn't be negative!

$$|\mathbf{u}| \geq 0$$


- And probably it should be zero only for the zero vector:

$$|\mathbf{u}| = 0 \iff \mathbf{u} = \mathbf{0}$$




Natural Properties of Length, Continued

- Also, if we scale a vector by a factor c , its norm will of course scale by the same amount:



- Finally, we know that the shortest path between two points is always along a straight line:



- (This final property is sometimes called the “pentagon inequality,” since the diagram looks like a pentagon.)


Norm—Formal Definition

- A norm is any function that assigns a number to each vector and satisfies the following properties for all vectors \mathbf{u} , \mathbf{v} , and all scalars a :
 - $|\mathbf{v}| \geq 0$
 - $|\mathbf{v}| = 0 \iff \mathbf{v} = \mathbf{0}$
 - $|a\mathbf{v}| = a|\mathbf{v}|$
 - $|\mathbf{u}| + |\mathbf{v}| \geq |\mathbf{u} + \mathbf{v}|$
- As before, you should be able to associate each of these properties with a concrete, geometric interpretation. (Visualize in your head the picture associated with each one.)

Euclidean Norm in Cartesian Coordinates

- A standard norm is the so-called *Euclidean norm* of n -vectors:

$$|\mathbf{u}| = |(u_1, \dots, u_n)| := \sqrt{\sum_{i=1}^n u_i^2}$$



Example: $\mathbf{u} = (4, 2)$

$$\begin{aligned} |\mathbf{u}| &= \sqrt{4^2 + 2^2} \\ &= 2\sqrt{5} \end{aligned}$$



Q: Does this formula satisfy all the natural, geometric properties of a norm?
(Answer in the slide comments!)

L² Norm of Functions

- Less familiar idea, but same basic intuition: the so-called *L² norm* measures the total magnitude of a function.
- Consider real-valued functions on the unit interval [0,1] whose square has a well-defined integral. The L² norm is defined as:

$$||f|| := \sqrt{\int_0^1 f(x)^2 dx}$$

- Not too different from the Euclidean norm, actually: just replaced a sum with an integral (which is kind of like a sum...).



Q: Does the formula above *exactly* satisfy all our desired properties for a norm?

L² Norm of Functions—Example


- Consider the function $f(x) := x\sqrt{3}$, defined over the unit interval $[0,1]$.
- $$||f|| := \sqrt{\int_0^1 f(x)^2 dx}$$

- Q: What is its L² norm?

- A:

$$||f||^2 = \int_0^1 3x^2 dx = \left[x^3 \right]_0^1 = 1^3 - 0^3 = 1$$




$$\Rightarrow ||f|| = \sqrt{1} = 1.$$



Note: for clarity we will use $|| \cdot ||$ for the norm of a function, and $| \cdot |$ for the norm of a vector in \mathbb{R}^n .

Inner Product—Motivation


- What else can we measure? In addition to magnitude, said that vectors have *orientation*. Just as norm measured length, **inner product** measures how well vectors “line up.”



Inner Product—Symmetry

- Will write inner product (also sometimes called the **scalar product** or **dot product**) using the notation $\langle u, v \rangle$ (some folks also write it as $u \cdot v$).
- When measuring the alignment of two vectors u, v , what are some natural properties you might expect?
- One “obvious” property: order shouldn’t matter, since u is just as well-aligned with v as v is with u :


$$\langle u, v \rangle = \langle v, u \rangle$$



- Equivalently, simply *re-naming* the vectors should have no effect on how well-aligned they are!


Inner Product—Projection & Scaling

- For unit vectors $|u|=|v|=1$, an inner product measures the extent of one vector along the direction of the other:




**Q: Is this property symmetric?
I.e., is the length of v along u the
same as the length of u along v ?**

- If we scale either of the vectors, the inner product also scales:



Inner Product—Positivity

- Also, a vector should always be aligned with itself, which we can express by saying that the inner product of a vector with itself should be positive (or at least, non-negative):



$$\langle \mathbf{u}, \mathbf{u} \rangle \geq 0$$

- In fact, if we continue to think of the inner product of a vector as the length of one vector along another then for unit-length vectors we must have

$$\langle \mathbf{u}, \mathbf{u} \rangle = 1$$

- Q: In general, then, what must $\langle \mathbf{u}, \mathbf{u} \rangle$ be equal to?
- A: Letting $\hat{\mathbf{u}} := \mathbf{u} / |\mathbf{u}|$, we have

$$\langle \mathbf{u}, \mathbf{u} \rangle = \langle |\mathbf{u}| \hat{\mathbf{u}}, |\mathbf{u}| \hat{\mathbf{u}} \rangle = |\mathbf{u}|^2 \langle \hat{\mathbf{u}}, \hat{\mathbf{u}} \rangle = |\mathbf{u}|^2 \cdot 1 = |\mathbf{u}|^2$$

Inner Product—Formal Definition

- An inner product is any function that assigns to any two vectors \mathbf{u}, \mathbf{v} a number $\langle \mathbf{u}, \mathbf{v} \rangle$ satisfying the following properties:


- $\langle \mathbf{u}, \mathbf{v} \rangle = \langle \mathbf{v}, \mathbf{u} \rangle$
- $\langle \mathbf{u}, \mathbf{u} \rangle \geq 0$
- $\langle \mathbf{u}, \mathbf{u} \rangle = 0 \iff \mathbf{u} = \mathbf{0}$
- $\langle a\mathbf{u}, \mathbf{v} \rangle = a\langle \mathbf{u}, \mathbf{v} \rangle$
- $\langle \mathbf{u} + \mathbf{v}, \mathbf{w} \rangle = \langle \mathbf{u}, \mathbf{w} \rangle + \langle \mathbf{v}, \mathbf{w} \rangle$

- Q: Which of these properties didn't we talk about? Can you argue that they make sense geometrically? (Discuss online!)

Inner Product in Cartesian Coordinates

- A standard inner product is the so-called Euclidean inner product, which operates on a pair of n -vectors:

$$\langle \mathbf{u}, \mathbf{v} \rangle = \langle (u_1, \dots, u_n), (v_1, \dots, v_n) \rangle := \sum_{i=1}^n u_i v_i$$



Example:

$$\mathbf{u} = (4, 1)$$

$$\mathbf{v} = (1, 3)$$

$$\langle \mathbf{u}, \mathbf{v} \rangle = 4 \cdot 1 + 1 \cdot 3 = 7$$




L² Inner Product of Functions—Example

- Just like we had a norm for functions, we can also define an inner product that measures how well two functions “line up”.
- E.g., for square-integrable functions on the unit interval:

$$\langle\langle f, g \rangle\rangle := \int_0^1 f(x)g(x) \, dx$$

Example: $f(x) := x^2$, $g(x) := (1 - x)^2$

$$\langle\langle f, g \rangle\rangle = \int_0^1 x^2(1 - x)^2 \, dx = \dots = \frac{1}{30}$$



**small number;
functions don't
“line up” much!**

Measuring Images, Other Signals?

- Many ways to measure “how big” a signal is (norm) or “how well-aligned” two signals are (inner product).
- Choice depends on application.
- Often, looks just like L^2 : integrate square over the domain.
- Could look quite different—for instance, measure norm of derivative (e.g., if interested in the *edges* of an image).




Linear Maps

- At the beginning, said linear algebra was study of **vector spaces** and **linear maps** between them.
- Have a pretty good handle on vector (and inner product) spaces.
- But what's a linear map? And why is it useful for graphics?
- We'll get to the 1st question in a moment. As for the 2nd question, a few reasons:
 - Computationally, easy to solve equations involving linear maps.
 - Basic transformations (rotation, translation, scaling) can be expressed as linear maps. *(Will see this in a later lecture!)*
 - Over a short distance, or a small amount of time, *all* maps can be approximated as linear maps. (Taylor's theorem). This approximation is used all over geometry, animation, rendering, image processing...

Linear Maps—Geometric Definition

- What is a linear map?
- Especially in graphics, can think about them *visually*.
- Example:



Key idea: *linear maps take lines to lines**

***...while keeping the origin fixed.**


Linear Maps—Algebraic Definition

- A map f is **linear** if it maps *vectors to vectors*, and if for all vectors \mathbf{u}, \mathbf{v} and scalars a we have

$$f(\mathbf{u} + \mathbf{v}) = f(\mathbf{u}) + f(\mathbf{v})$$

$$f(a\mathbf{u}) = af(\mathbf{u})$$

- In other words: if it doesn't matter whether we add the vectors and then apply the map, or apply the map and then add the vectors (ditto for scaling):




Linear Maps—Coordinate Definition

- For maps between \mathbb{R}^m and \mathbb{R}^n (e.g., a map from 2D to 3D), we can give an even more explicit definition.
- A map is linear if it can be expressed as

$$f(u_1, \dots, u_m) = \sum_{i=1}^m u_i \mathbf{a}_i$$


- In other words, if it is a linear combination of a fixed set of vectors \mathbf{a}_i :



Q: Is $f(x) := ax + b$ a linear function?

Linear vs. Affine Maps

- No! But it's easy to be fooled, since the graph looks like a line:



- However, it's not a line through the *origin*, i.e., $f(0) \neq 0$.
- Another way to see it's not linear? Doesn't preserve sums:

$$\begin{aligned} f(x_1 + x_2) &= a(x_1 + x_2) + b = ax_1 + ax_2 + b \\ f(x_1) + f(x_2) &= (ax_1 + b) + (ax_2 + b) = ax_1 + ax_2 + 2b \end{aligned}$$

- Yet another way: not (just) a linear combination of fixed a's.
- Important computer graphics magic trick: turn affine transformations (e.g., translation) into linear ones via *projective coordinates* (will see later on!)

More interesting question:

Q: Is $f(u) := \int_0^1 u(x) dx$ a linear map?

**(Think about it—it will be
part of your homework!)**

Span


Q: Geometrically, what is the *span* of two vectors \mathbf{u} , \mathbf{v} ?

A: The span is the set of all vectors that can be written as a linear combination of \mathbf{u} and \mathbf{v} , i.e., vectors of the form

$$a\mathbf{u} + b\mathbf{v}$$


for any two numbers a , b .

More generally: $\text{span}(\mathbf{u}_1, \dots, \mathbf{u}_k) = \left\{ \mathbf{x} \in V \mid \mathbf{x} = \sum_{i=1}^k a_i \mathbf{u}_i, a_1, \dots, a_k \in \mathbb{R} \right\}$



Span & Linear Maps

- Just a bit of language—can connect “span” and “linear map”:
- *“The **image** of any linear map is the span of some collection of vectors.”*




Q: What's the **image of a function?**

Basis

- Span is also closely related to the idea of a *basis*.
- In particular, if we have exactly n vectors $\mathbf{e}_1, \dots, \mathbf{e}_n$ such that

$$\text{span}(\mathbf{e}_1, \dots, \mathbf{e}_n) = \mathbb{R}^n$$

- Then we say that these vectors are a **basis** for \mathbb{R}^n .
- Note: many different choices of basis!
- Q: Which of the following are bases for the 2D plane ($n=2$)?




Orthonormal Basis

- Most often, it is convenient to have basis vectors that are (i) unit length and (ii) mutually orthogonal.
- In other words, if $\mathbf{e}_1, \dots, \mathbf{e}_n$ are our basis vectors then

$$\langle \mathbf{e}_i, \mathbf{e}_j \rangle = \begin{cases} 1, & i = j \\ 0, & \text{otherwise.} \end{cases}$$


- This way, the geometric meaning of the sum $u_1^2 + \dots + u_n^2$ is maintained: it is the length of the vector \mathbf{u} .



Common bug: projecting a vector onto a basis that is NOT orthonormal while continuing to use standard norm / inner product.

Gram-Schmidt

- Given a collection of basis vectors $\mathbf{a}_1, \dots, \mathbf{a}_n$, how do we find an **orthonormal** basis $\mathbf{e}_1, \dots, \mathbf{e}_n$?
- Gram-Schmidt algorithm (or “process”):
 - normalize the first vector (i.e., divide by its length)
 - subtract any component of the 1st vector from the 2nd one
 - normalize the 2nd one
 - subtract any component of the first two from the 3rd one
 - ...



***WARNING: for large number of vectors / nearly parallel vectors, not the best algorithm...**

Gram-Schmidt—Example

- Common task in graphics: have a triangle in 3D, need orthonormal basis for the plane containing the triangle
- Strategy: apply Gram-Schmidt to (any) pair of edge vectors


$$\mathbf{u} := p_1 - p_0$$

$$\mathbf{v} := p_2 - p_0$$

$$\mathbf{e}_1 := \mathbf{u} / |\mathbf{u}|$$

$$\tilde{\mathbf{v}} := \mathbf{v} - \langle \mathbf{v}, \mathbf{e}_1 \rangle \mathbf{e}_1$$


$$\mathbf{e}_2 := \tilde{\mathbf{v}} / |\tilde{\mathbf{v}}|$$





Food for thought: in 3D, does the *order* of bases matter? I.e., $(\mathbf{e}_1, \mathbf{e}_2)$ vs. $(\mathbf{e}_2, \mathbf{e}_1)$?

Orthonormal Basis for Functions

- Functions are also vectors. Can we also decompose them into orthogonal “components?”
- Surprisingly, the answer is “YES!”
- Simple example: square-integrable 2π -periodic functions (i.e., $f(x) = f(x+2k\pi)$ for all integers k)




- Orthogonal basis given by sinusoids: $\cos(nx)$, $\sin(mx)$, $m, n \in \mathbb{N}$
- Easy to normalize, but why are these bases *orthogonal*?



Projection of Function onto Sinusoids


- Projection onto bases works just like vectors in \mathbb{R}^n :




$$c_k := \langle\langle f(x), \cos(kx) \rangle\rangle$$

$$s_k := \langle\langle f(x), \sin(kx) \rangle\rangle$$

$$\Rightarrow f(x) = \sum_{k \in \mathbb{N}} c_k \cos(kx) + s_k \sin(kx)$$



- Decomposes signal into “frequencies”:




lots of low- and mid-frequency oscillation


not as much high-frequency oscillation

Frequency Decomposition of Signals


- More generally, this idea of projecting a signal onto different “frequencies” is known as **Fourier decomposition**
- Can be applied to all sorts of signals; basic tool used across, image processing, rendering, geometry, physical simulation...
- Will have plenty more to say as course goes on!




$m = 40$









$m=200$



$m=500$



$m=1000$



Systems of Linear Equations

- **Extremely common task in computer graphics algorithms: solve large systems of linear equations.**
- **Represent everything from light bouncing around a scene, to time evolution of physical phenomena (fluids, tissue, rubber...), to basic geometric problems (e.g., making a flat map of a curved surface).**
 - **Typically, very “sparse” systems, meaning many, many variables, but each variable appears in only a small number of equations.**
 - **Often the computational bottleneck in graphics algorithms.**
 - **Only last ~10 years or so people even willing to solve such big systems.**
 - **Use specialized libraries for sparse numerical linear algebra.**

System of Linear Equations

- A system of linear equations is exactly what it sounds like: a bunch of equations where left-hand side is a linear function, right hand side is constant. E.g.,

$$\begin{array}{rcl} x + 2y & = & 3 \\ 4x + 5y & = & 6 \end{array}$$

- Unknown values are sometimes called “degrees of freedom” (DOFs); equations sometimes called “constraints”
- Goal: solve for DOFs that simultaneously satisfy constraints:

$$\begin{array}{l} x = 3 - 2y \\ 4(3 - 2y) + 5y = 6 \end{array}$$


$$\boxed{\begin{array}{l} y = 2 \\ x = -1 \end{array}}$$

What does solving a linear system *mean*?


Linear System, Visualized

- Of course, a linear system can be used to represent many different practical tasks (simulation, processing, etc.).
- But for any linear system, there are some good mental models to visualize:

Find the point where two lines meet:






GIVEN a point b , FIND the point x that maps to it:




Uniqueness, Existence of Solutions

- Of course, not all linear systems can be solved! (And even those that can be solved may not have a unique solution.)



Bilinear & Quadratic Forms

- In homework, you will take a look at *bilinear and quadratic forms*.
- Core idea in graphics, since many modern algorithms are formulated in terms of *energy minimization*, often involving quadratic form.
- Already familiar with one example: inner product / norm!



Wait, what about matrices?!

Matrices in Linear Algebra

- Linear algebra often taught from the perspective of *matrices*, i.e., pushing around little blocks of numbers.
- But linear algebra is not fundamentally *about* matrices.
- As you've just seen, you can understand almost all the basic concepts without ever touching a matrix!
- Likewise, matrices can interfere with understanding / lead to confusion, since the same object (a block of numbers) is used to represent many different things (linear map, quadratic form, ...) in many different bases.
- Still, VERY useful!
 - symbolic manipulation
 - numerical computation


$$\begin{bmatrix} 1 & 7 & 3 \\ 4 & 9 & 2 \\ 0 & 1 & 1 \end{bmatrix}$$

What does this thing mean/
encode/do/represent?

Representing Linear Maps via Matrices

- Key example: suppose I have a linear map

$$f(\mathbf{u}) = u_1 \mathbf{a}_1 + u_2 \mathbf{a}_2$$



- How do I encode as a matrix?
- Easy: “a” vectors become matrix columns:

$$A := \begin{bmatrix} a_{1,x} & a_{2,x} \\ a_{1,y} & a_{2,y} \\ a_{1,z} & a_{2,z} \end{bmatrix}$$

- Now, matrix-vector multiply recovers original map:

$$\begin{bmatrix} a_{1,x} & a_{2,x} \\ a_{1,y} & a_{2,y} \\ a_{1,z} & a_{2,z} \end{bmatrix} \begin{bmatrix} u_1 \\ u_2 \end{bmatrix} = \begin{bmatrix} a_{1,x}u_1 + a_{2,x}u_2 \\ a_{1,y}u_1 + a_{2,y}u_2 \\ a_{1,z}u_1 + a_{2,z}u_2 \end{bmatrix} = u_1 \mathbf{a}_1 + u_2 \mathbf{a}_2$$

**Don't worry: if you love matrices, there will
be more of them in your homework!**

QUIZ[1]

Since you already have plenty of exercises related to today's material (in your homework), today's quiz is to

Post a question or comment on the course slides.

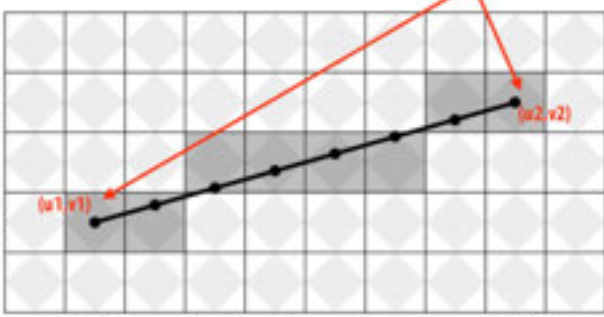
Two-word comments like “Cool, man!” will not receive any credit—remember that these quizzes are graded based on effort alone. (So make an effort!)

To keep you coming to class (and to make grading easier...), please *print out the slide you commented on and bring it to the next lecture.*

Incremental line rasterization

- Let's say a line is represented with integer endpoints: $(u_1, v_1), (u_2, v_2)$
- Slope of line: $s = (v_2 - v_1) / (u_2 - u_1)$
- Consider a very easy special case:
 - $u_1 < u_2, v_1 < v_2$ (line points toward upper-right)
 - $0 < s < 1$ (more change in x than y)

Assume integer coordinates are at pixel centers



```
v = v1;
for( u=u1; u<=u2; u++ )
{
    v += s;
    draw( u, round(v) );
}
```

Common optimization: rewrite algorithm to use only integer arithmetic (Bresenham algorithm)

CMU 15-462/662, Fall 2015

shhhh a day ago

Can't you just tweak the algorithm to work in parallel? Something along the lines of: for (i = 0; i < u2-u1; i++) { draw (u1+i, round(v1+(i*s))) }

Prompt Edit Delete Archive [0 Upvote Downvote]

keenana a day ago

Are there any methods that can break the line up into sections (especially for high resolution outputs)?

Can't you just tweak the algorithm to work in parallel? Something along the lines of: for (i = 0; i < u2-u1; i++) { draw (u1+i, round(v1+(i*s))) }


Yes, in fact a modern way to render lines (strange as it may seem at first) is to render a quadrilateral, which itself is actually drawn by the graphics hardware as a pair of triangles—imagine splitting the quad along its diagonal. From there, the triangles are efficiently rendered in parallel via an algorithm like the ones we'll look at when we talk about rasterization.

Of course that leaves the question: why do all the work of rendering a line as two triangles, when it seems much simpler to have a specialized algorithm for lines? The answer is: because if you can decompose all of your draw operations into one atomic operation (in this case, drawing triangles), then you can make the hardware much simpler. And from there, you can really make things fly (e.g., by duplicating the units that draw triangles many times, and optimizing the heck out of them). If you need specialized algorithms for each type of primitive, you have to share space on the chip with lots of different algorithms, many of which may not be used much (lines are far less common than triangles in modern applications).

Next time: Math (P)Review Part II


- Eigenvalue problems
- Vector calculus
- Complex numbers

ϕ




$\text{grad } \phi$

X



$\text{div } X$

Y



$\text{curl } Y$